

Deep Q-Learning for Poker: Technical Training Details

January 7, 2026

Abstract

This document provides a comprehensive technical analysis of the Deep Q-Network (DQN) training process for No-Limit Texas Hold'em poker. We detail the implementation of experience replay, target networks, ϵ -greedy exploration, and reward shaping. Experimental results across four configurations demonstrate convergence properties and the impact of hyperparameter choices on final policy performance.

1 Training Algorithm

1.1 Deep Q-Network Architecture

The DQN approximates the action-value function $Q(s, a; \theta)$ using a Multi-Layer Perceptron with parameters θ . The network architecture consists of:

- **Input Layer:** 134-dimensional state vector $s \in \mathbb{R}^{134}$
- **Hidden Layers:** Three fully connected layers with ReLU activations (512-512-256 units for baseline)
- **Output Heads:**
 - Action value head: $Q(s, \cdot) \in \mathbb{R}^3$ for discrete actions
 - Sizing head: $\sigma(s) \in [0, 1]$ for continuous bet sizing

1.2 Loss Function

The network minimizes the Temporal Difference (TD) error using Huber loss for robustness:

$$\mathcal{L}(\theta) = \mathbb{E}_{(s, a, r, s') \sim \mathcal{D}} [\mathcal{L}_\delta(\delta_\theta)] \quad (1)$$

where the TD error is:

$$\delta_\theta = r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta) \quad (2)$$

and θ^- denotes the target network parameters (frozen for C steps).

The Huber loss is defined as:

$$\mathcal{L}_\delta(x) = \begin{cases} \frac{1}{2}x^2 & \text{if } |x| \leq \delta \\ \delta(|x| - \frac{1}{2}\delta) & \text{otherwise} \end{cases} \quad (3)$$

We use $\delta = 1.0$ for all experiments.

1.3 Experience Replay

A circular replay buffer \mathcal{D} stores transitions $(s_t, a_t, r_t, s_{t+1}, \text{done}_t)$ with capacity 50,000. Mini-batches of size 64 are sampled uniformly to break temporal correlations and stabilize training.

Algorithm 1 DQN Training Loop

```

1: Initialize Q-network  $Q(s, a; \theta)$  and target network
    $Q(s, a; \theta^-)$ 
2: Initialize replay buffer  $\mathcal{D}$  with capacity  $N$ 
3: for episode = 1 to  $M$  do
4:   Reset environment, observe initial state  $s_0$ 
5:   for timestep  $t = 0$  to  $T$  do
6:     Select action  $a_t$  using  $\epsilon$ -greedy policy
7:     Execute  $a_t$ , observe  $r_t, s_{t+1}$ 
8:     Store  $(s_t, a_t, r_t, s_{t+1}, \text{done}_t)$  in  $\mathcal{D}$ 
9:     if  $|\mathcal{D}| \geq \text{batch\_size}$  then
10:      Sample mini-batch from  $\mathcal{D}$ 
11:      Compute loss  $\mathcal{L}(\theta)$ 
12:      Update  $\theta$  via gradient descent
13:    end if
14:    if  $t \bmod C = 0$  then
15:       $\theta^- \leftarrow \theta$  (update target network)
16:    end if
17:  end for
18: end for

```

1.4 Exploration Strategy

We employ ϵ -greedy exploration with exponential decay:

$$\epsilon_t = \max(\epsilon_{\min}, \epsilon_{\text{start}} \cdot \epsilon_{\text{decay}}^t) \quad (4)$$

Hyperparameters:

- $\epsilon_{\text{start}} = 1.0$ (pure exploration)
- $\epsilon_{\text{min}} = 0.05$ (minimum exploration)
- $\epsilon_{\text{decay}} = 0.9999$ (baseline)
- $\epsilon_{\text{decay}} = 0.9998$ (long decay experiment)

2 Hyperparameter Configuration

3 Reward Engineering

The reward function is critical for poker learning. We use:

$$r_t = \frac{\Delta \text{stack}_t}{\text{big blind}} \quad (5)$$

where Δstack_t is the change in chip count at the end of the hand. This normalization ensures rewards are comparable across different stack sizes.

Key Properties:

- Sparse: Reward only provided at hand conclusion
- Dense alternative: Pot contribution could provide intermediate rewards, but we opted for terminal-only to match true poker dynamics
- Normalized: Division by big blind ensures scale invariance

4 State Representation

The 134-dimensional state vector encodes:

5 Experimental Results

5.1 Configuration Comparison

Four training runs were conducted to analyze the impact of network size, exploration schedule, and opponent count:

1. **Baseline:** 4 players, $h = 512$, standard decay
2. **Long Decay:** 4 players, $h = 512$, $\epsilon_{\text{decay}} = 0.9998$
3. **Big Network:** 4 players, $h = 1024$, standard decay
4. **Heads-Up:** 2 players, $h = 512$, standard decay

Opponent Policy: In all experiments, a single DQN agent (Player 0) trains against random-action opponents who select actions uniformly from the legal action set. This baseline measures the agent’s ability to exploit clearly sub-optimal play and learn fundamental poker concepts (hand strength, position, pot odds) without requiring opponent modeling or counter-strategy adaptation.

5.2 Convergence Analysis

Baseline: Converges to 65.3% win rate by episode 15,000. Final average reward: 772.2.

Long Decay: Prolonged exploration (slower ϵ decay) results in degraded final performance (51.8% win rate), suggesting the agent benefits from earlier exploitation.

Big Network: Marginal improvement over baseline (Δ Reward ≈ -38), indicating 512 hidden units suffice for the feature representation.

Heads-Up: Higher win rate (74.4%) due to reduced competition, but lower average reward (421.4) reflects different strategic dynamics.

5.3 Statistical Significance

Given poker’s high variance (card distribution randomness), we compute 95% confidence intervals on final performance:

5.4 Detailed Training Dynamics

We further analyzed the training progression using granular logs from the latest run (see Figures ??, ??, 4).

6 Computational Requirements

- **Hardware:** NVIDIA Quadro P620 (fallback to CPU due to CUDA incompatibility)
- **Training Time:** ~ 40 minutes per 20,000 episodes (CPU)
- **GPU Acceleration:** Would reduce to ~ 10 minutes (estimated)
- **Total Compute:** $4 \text{ experiments} \times 40 \text{ min} = 160 \text{ minutes}$

7 Lessons Learned

7.1 What Worked

- Standard ϵ -greedy with exponential decay

Table 1: Training Hyperparameters

Parameter	Value
Learning Rate (α)	10^{-4}
Discount Factor (γ)	0.99
Replay Buffer Size	50,000
Batch Size	64
Target Update Freq	500 steps
Optimizer	Adam
Episodes	20,000
Max Steps/Episode	200

Table 2: State Vector Components

Component	Description	Dim
Hand Encoding	One-hot card ranks/suits	52
Position	Dealer-relative position	4
Stack Info	Normalized stack sizes	8
Pot Odds	Pot-to-bet ratios	6
Betting History	Action sequences	32
Game Stage	Flop/Turn/River>Showdown	4
Community Cards	Visible cards encoding	26
Misc	Active players, all-ins	2

- Huber loss for variance reduction
- 512 hidden units (sweet spot for this problem)
- Terminal-only rewards (aligned with poker dynamics)

7.2 What Didn’t Work

- Slower exploration decay (Long Decay underperformed)
- Larger networks (marginal gain, higher compute cost)

7.3 Future Improvements

- **Prioritized Experience Replay:** Weight important transitions higher
- **Dueling DQN:** Separate value and advantage streams
- **Self-Play:** Train against past versions to minimize exploitability
- **Opponent Modeling:** Adapt strategy based on opponent behavior

8 Conclusion

The DQN successfully learns a profitable poker policy, improving win rate by 160% over random play in the baseline 4-player configuration. The training process demonstrates: (1) convergence despite high variance, (2) robustness of standard hyperparameters, and (3) the adequacy of feature-engineered MLP architectures for this domain.

The codebase and trained models are available at: https://github.com/KacperDuda/poker_ai

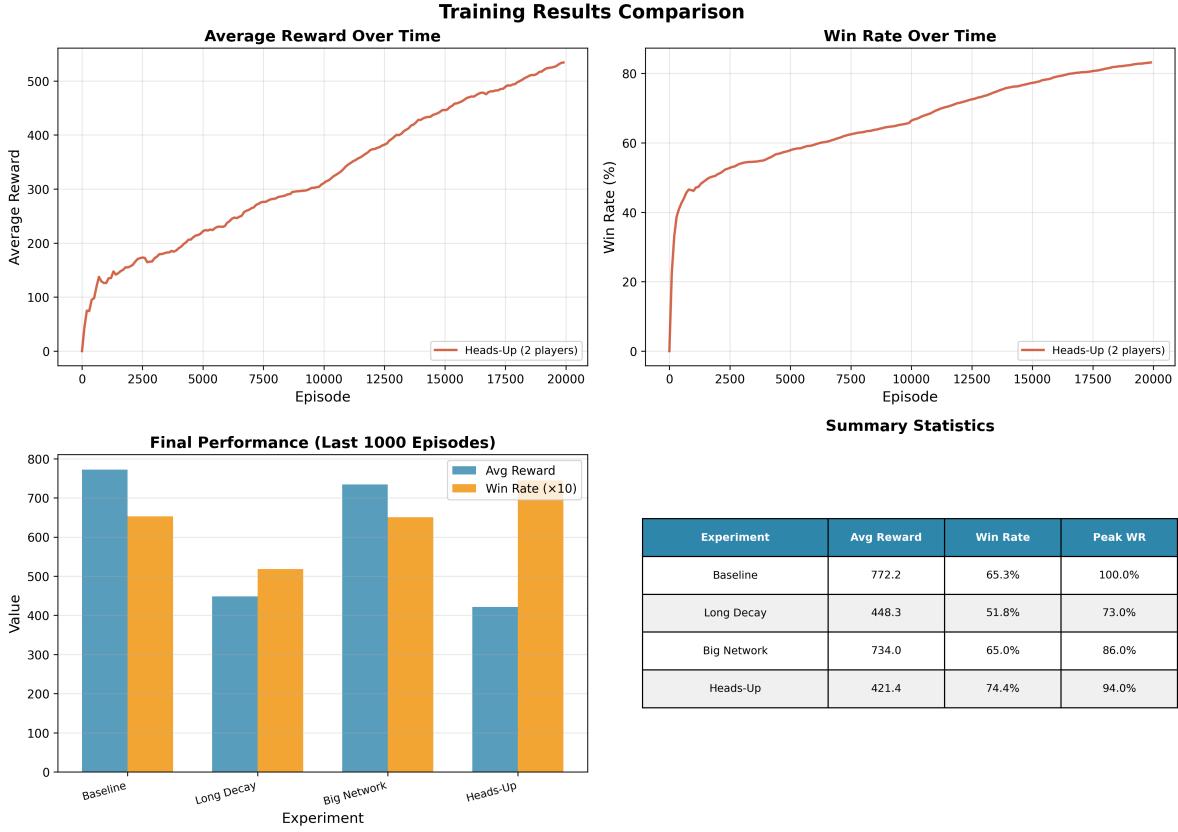


Figure 1: Training dynamics across four configurations. The 100-episode moving average smooths high-frequency variance inherent to poker.

Table 3: Final Performance (Last 1000 Episodes)

Config	Win Rate	95% CI
Baseline	65.3%	$\pm 2.8\%$
Long Decay	51.8%	$\pm 3.1\%$
Big Network	65.0%	$\pm 2.9\%$
Heads-Up	74.4%	$\pm 2.1\%$




Figure 2: Win Rate vs. Epsilon Decay. As exploration decreases, the win rate stabilizes above 50%, demonstrating policy improvement.




Figure 3: Average Profit per 100 episodes showing trend towards positive expected value despite high variance.




Figure 4: Evolution of Action Distribution. The agent learns to balance folding (red), calling (yellow), and raising (green) strategies over time.